

419 lines (284 loc) · 29.9 KB

Day 21 Task: Docker Essential Concepts

PARAG PALLAV SINGH

Docker Essentials

Interview Questions

Docker Interview

Docker is a good topic to ask in DevOps Engineer Interviews, mostly for freshers. One must surely try these questions in order to be better in Docker

Questions

1. What is the Difference between an Image, Container and Engine?

Aspect	Image	Container	Engine
Definition	Static snapshot of an application	Running instance of an image	Software that manages and runs containers
Contents	Includes code, runtime, libraries, etc.	Includes the image and runtime context	Provides tools and services for containers
Portability	Highly portable	Portable	Enables portability and orchestration
Isolation	Not executable on its own	Isolated runtime environment	Manages isolation and resource allocation

Aspect	Image	Container	Engine
Lifecycle	Created, stored, and shared	Created, started, stopped, and removed	Installed, configured, and operated
Purpose	Basis for creating containers	Runnable instance of an application	Manages containerization and deployment
Interaction	Used to create containers	Created from images, runs applications	Used to build, manage, and run containers
Persistence	Can be stored in a registry	Exists only while running	Persistent management of containers
Networking	No network configuration	Can have network settings	Manages networking for containers
Storage	No state storage	Can have state stored inside	Coordinates storage resources
Example Technologies	Docker images, OCI images	Docker containers, Kubernetes pods	Docker Engine, containerd, Kubernetes

2. What is the Difference between the Docker command COPY vs ADD?

Aspect	COPY	ADD
Purpose	Copies files and directories	Copies files and directories
Source	Can copy from the build context	Can copy from the build context
Destination	Only works with the container's file system	Can also extract archives and handle URLs
Extraction	Does not automatically extract archives	Can automatically extract common archives (e.g., <code>.tar</code> , <code>.gzip</code>)
URL Handling	Cannot handle URLs	Can download from URLs
Ownership & Permissions	Preserves ownership and permissions	Copies with default permissions
Cache Busting	Only invalidates cache if source file changes	Invalidates cache if source URL or contents change
Use Cases	Best for simple file copying	Offers additional features like URL handling and archive extraction

3. What is the Difference between the Docker command CMD vs RUN?

Aspect	CMD	RUN
Purpose	Specifies the default command for the container	Executes commands during image build
Usage	Typically used once at the end of a Dockerfile	Can be used multiple times in a Dockerfile
Command Execution	Executes only when the container starts	Executes during image build
Overriding	Can be overridden by a command during container run	Commands are fixed in the image
Default	Provides the default behavior for the container	Does not change the container behavior
Parameters	Executable command and optional arguments	Shell command(s) and arguments to execute
Example	<code>CMD ["nginx", "-g", "daemon off;"]</code>	<code>RUN apt-get update && apt-get install -y curl</code>

4. How Will you reduce the size of the Docker image?

Reducing the size of a Docker image is important for optimizing storage, transfer times, and resource consumption. Here are several strategies you can employ to achieve a smaller Docker image size:

- 1. Use Minimal Base Images:** Choose lightweight base images like Alpine Linux or BusyBox instead of full-fledged operating systems. These minimal images contain only essential components, leading to smaller image sizes.
- 2. Multi-Stage Builds:** Utilize multi-stage builds to separate build-time dependencies from runtime artifacts. This involves creating multiple build stages in the Dockerfile, where the final stage only includes the necessary runtime components. This can significantly reduce the image size.
- 3. Minimize Layers:** Each instruction in a Dockerfile creates a new layer in the image. Minimize the number of instructions to reduce the number of layers. Combine multiple commands into a single RUN instruction when possible.
- 4. Use COPY Instead of ADD:** Use the `COPY` command instead of `ADD` to copy files into the image. `COPY` is simpler and doesn't perform any automatic extraction, reducing the chances of unintended increase in image size.

5. **Cleanup Unnecessary Files:** Remove temporary files, caches, and any files that are not required in the final image. Use the `RUN` instruction to delete files immediately after they've been used to avoid unnecessary layers.
6. **Avoid Installing Unnecessary Packages:** When installing software packages, only include the dependencies required for your application to run. Minimize the number of extra packages that get installed.
7. **Use Specific Version Tags:** When installing software, specify exact version numbers for packages and dependencies. Avoid using generic or latest tags as they might result in larger images when dependencies change.
8. **Compress Files Before Adding:** If you need to include large files, compress them before adding them to the image. Then, you can decompress them during runtime.
9. **Use .dockerignore:** Create a `.dockerignore` file in your project directory to exclude unnecessary files and directories from being added to the image during the build process.
10. **Alpine Package Optimization:** When using Alpine Linux, you can optimize package installation by specifying the `--no-cache` flag in the `RUN` command to prevent the creation of cache files.
11. **Optimize Dockerfile Instructions:** Optimize the order of Dockerfile instructions. Place instructions that change infrequently (e.g., installing dependencies) before instructions that change frequently (e.g., copying source code).
12. **Clean Up Apt and Yum Cache:** If using Debian-based (`apt`) or Red Hat-based (`yum`) package managers, clean up the package cache in the same `RUN` command after installing packages.

5. Why and when to use Docker?

1. **Isolation and Consistency:** Docker provides containerization, isolating applications and their dependencies from the underlying system. This ensures consistent behavior across various environments, from development to production, reducing the "it works on my machine" problem.
2. **Portability:** Docker containers encapsulate the application and its dependencies into a single package. This makes it easy to move applications between different hosts, cloud platforms, and on-premises environments, fostering a "write once, run anywhere" approach.
3. **Resource Efficiency:** Docker containers share the host OS kernel, consuming fewer resources compared to traditional virtualization. This allows you to run more applications on the same infrastructure, optimizing resource utilization.
4. **Rapid Deployment and Scaling:** Docker's lightweight nature enables fast application deployment and scaling. You can easily spin up new containers or replicate existing ones to handle varying levels of load, improving application responsiveness.

5. **Microservices and DevOps:** Docker is well-suited for microservices architectures and DevOps practices. It allows teams to develop, test, and deploy small, independent services in parallel, enabling quicker iterations, continuous integration, and continuous deployment.

6. Explain the Docker components and how they interact with each other.

Docker consists of several key components that work together to enable containerization, deployment, and management of applications. Here's a brief overview of these components and how they interact:

1. **Docker Daemon:** The Docker daemon (dockerd) is a background process running on the host system. It's responsible for building, running, and managing containers. It listens for Docker API requests and communicates with the container runtime to manage containers' execution.
2. **Docker Client:** The Docker client (docker) is the command-line interface used by users to interact with the Docker daemon. It allows users to issue commands to build, run, stop, and manage containers. The client communicates with the Docker daemon via the Docker API.
3. **Docker Images:** Docker images are the building blocks of containers. An image is a lightweight, standalone, executable software package that includes the application code, runtime, libraries, and dependencies. Images are created from a set of instructions defined in a Dockerfile and can be stored in registries for distribution.
4. **Docker Containers:** Containers are instances of Docker images. They encapsulate the application, its dependencies, and runtime settings in an isolated environment. Containers share the host OS kernel but have their own isolated file system, networking, and process space.
5. **Docker Compose:** Docker Compose is a tool for defining and running multi-container Docker applications. It uses a YAML file to define services, networks, and volumes for different containers in a single application stack. Compose simplifies managing complex applications with multiple interconnected services.
6. **Docker Registry:** Docker registries are repositories for storing and distributing Docker images. The most common public registry is Docker Hub, but private registries can also be set up for secure image storage and distribution within an organization.
7. **Docker Network:** Docker networks allow containers to communicate with each other and with the outside world. They provide isolation and manage networking resources for containers, enabling seamless communication between containers on the same or different hosts.
8. **Docker Volumes:** Docker volumes are a way to manage persistent data for containers. Volumes can be mounted into containers, allowing data to be stored separately from the container's file system and persist even after a container is removed.

7. Explain the terminology: Docker Compose, Docker File, Docker Image, Docker Container?

1. **Docker Compose:** Docker Compose is a tool for defining and managing multi-container Docker applications. It uses a YAML file to define the services, networks, and volumes that make up an application stack. Compose simplifies the process of orchestrating and running interconnected containers by allowing you to define and manage the entire application environment in a single configuration file.
2. **Dockerfile:** A Dockerfile is a text file that contains a set of instructions for building a Docker image. It defines the base image, sets up the environment, installs dependencies, copies files, and configures settings needed for your application. Dockerfiles are used to automate the creation of consistent and reproducible Docker images.
3. **Docker Image:** A Docker image is a standalone, executable software package that includes everything needed to run an application, such as the code, runtime, libraries, and dependencies. Images are created from Dockerfiles and can be shared, stored in registries, and used to create containers. Images serve as templates for containers and ensure consistent application behavior across different environments.
4. **Docker Container:** A Docker container is a runnable instance of a Docker image. Containers are isolated environments that encapsulate an application and its dependencies. They share the host OS kernel but have their own isolated file system, networking, and process space. Containers are lightweight, portable, and can be started, stopped, and replicated quickly.

Together, these components enable efficient application development, deployment, and management through containerization.

8. In what real scenarios have you used Docker?

Here are some of the common scenarios where Docker is commonly used in the software development and deployment world:

1. **Microservices Architecture:** Docker is often used to containerize individual microservices within a larger application. This approach allows each microservice to be developed, tested, and deployed independently, improving scalability and maintainability.
2. **DevOps Practices:** Docker is a key tool in implementing DevOps practices. It helps in creating consistent development, testing, and production environments, reducing discrepancies between different stages of the software development lifecycle.
3. **Continuous Integration and Continuous Deployment (CI/CD):** Docker images can be easily integrated into CI/CD pipelines, ensuring that the same image used during development is also used for testing and deployment. This helps catch issues earlier in the development process.
4. **Local Development:** Developers often use Docker to replicate the production environment on their local machines. This ensures that what they build and test locally closely matches the production environment.

- 5. **Portability and Cloud Deployment:** Docker's portability is particularly useful when deploying applications to the cloud. It allows you to package an application and its dependencies into a single container, making it easy to move and run across different cloud providers.
- 6. **Isolation for Testing:** Docker containers provide isolation, making it easier to test software in a controlled environment without affecting the host system. This is especially valuable for testing different software configurations.
- 7. **Legacy Application Modernization:** Docker can help modernize legacy applications by packaging them in containers. This allows older applications to run on modern infrastructure and take advantage of containerization benefits.
- 8. **Resource Optimization:** By using Docker's lightweight containers, you can optimize resource utilization on servers, running multiple containers on a single host while maintaining isolation.
- 9. **Big Data and Analytics:** Docker is used in big data environments to create isolated containers for various analytics tools, databases, and frameworks, ensuring consistent environments for data processing and analysis.
- 10. **Edge Computing:** Docker's lightweight and consistent runtime makes it suitable for deploying applications at the edge of the network, where resource constraints and diverse hardware environments exist.

9. Docker vs Hypervisor?

Aspect	Docker (Containerization)	Hypervisor (Virtualization)
Isolation Level	Operating system-level virtualization	Full hardware-level virtualization
Overhead	Minimal resource overhead	Higher resource overhead
Performance	Near-native performance	Slightly lower performance due to overhead
Resource Usage	Shares host OS kernel, efficient resource usage	Emulates complete hardware, resource-intensive
Portability	Highly portable across environments	Less portable, requires more configuration
Start Time	Near-instant startup and shutdown	Longer startup and shutdown times
Density	Higher container density on a host	Lower VM density due to resource needs
Networking	Shared network stack with host and other containers	Separate virtual network stack for each VM

Aspect	Docker (Containerization)	Hypervisor (Virtualization)
Storage	Shared file system with host and other containers	Virtual disk images for each VM
Orchestration	Integrated with tools like Docker Compose and Kubernetes	Requires separate management tools
Use Cases	Microservices, DevOps, lightweight apps	Legacy applications, complex environments

10. What are the advantages and disadvantages of using docker?

Advantages	Disadvantages
Isolation: Containers provide application isolation without the overhead of full virtualization.	Learning Curve: New concepts and tools to master.
Portability: Consistent environments across development, testing, and production.	Resource Overhead: Slight performance overhead due to shared kernel.
Resource Efficiency: Containers share the host OS kernel, leading to efficient resource utilization.	Limited OS Support: Works best on Linux; Windows and macOS support has limitations.
Rapid Deployment: Containers start quickly and can be scaled up or down easily.	Security Concerns: Misconfiguration can lead to vulnerabilities.
Version Control: Images provide versioned snapshots of applications and dependencies.	Complex Networking: Networking between containers can be complex.
Microservices: Ideal for microservices architecture and modular application design.	Persistence: Containers are typically stateless, requiring separate solutions for data storage.
Ecosystem: Rich ecosystem with tools like Docker Compose and Kubernetes for orchestration.	Image Size: Images can be large, especially with full environments.
DevOps Integration: Facilitates CI/CD pipelines and DevOps practices.	Incompatibility: Some legacy applications might not work well in containers.
Isolation and Compatibility: Containers can run side by side with different dependencies.	Continuous Maintenance: Images and containers need regular updates and management.
Development Consistency: Replicate production environment on local machine for development.	Lack of GUI: Docker is primarily command-line driven; graphical interfaces are limited.

11. What is a Docker namespace?

In Docker, namespaces are a key component of containerization that provide process isolation and resource separation between containers and the host system. They create a virtualized environment for each container, isolating processes, filesystems, networks, and other resources. This allows multiple containers to run on the same host without interfering with each other or the host system.

Here's a brief overview of namespaces along with an example:

Namespace Types in Docker:

1. **PID Namespace (pid):** Provides process isolation. Each container sees only its own set of processes, with process IDs (PIDs) unique within the namespace.
2. **Network Namespace (net):** Isolates networking resources, including network interfaces, IP addresses, routing tables, and firewall rules.
3. **Mount Namespace (mnt):** Provides filesystem isolation. Containers have their own filesystem hierarchy, separate from the host and other containers.
4. **UTS Namespace (uts):** Isolates hostname and domain name, allowing each container to have its own identity for these attributes.
5. **IPC Namespace (ipc):** Isolates inter-process communication (IPC) resources such as shared memory segments and semaphores.
6. **User Namespace (user):** Isolates user and group IDs, allowing containers to have their own user and group mappings.

Example:

Let's consider a scenario where two Docker containers are running on the same host, both serving web applications. Each container needs to run its own web server on port 80.

In this case, the network namespace ensures that both containers can use port 80 without conflict. Each container's network namespace is isolated, allowing them to have their own IP addresses, network interfaces, and routing tables.

Container 1:

- Network Namespace: Isolated
- IP Address: 172.17.0.2
- Port 80: Mapped to 8080

Container 2:

- Network Namespace: Isolated
- IP Address: 172.17.0.3

- Port 80: Mapped to 8081

12. What is a Docker registry?

A Docker registry is a centralized repository that stores Docker images. It allows you to manage, distribute, and share Docker images across different environments and teams. Docker Hub is one of the most popular public Docker registries, but you can also set up private registries to control image distribution within your organization.

Here's a brief overview of a Docker registry along with an example:

Public Docker Registry (Docker Hub): Docker Hub is a public registry where you can find and share Docker images. It's commonly used for open-source projects and public applications. You can pull images from Docker Hub to your local machine or directly to your server for deployment.

Private Docker Registry: A private Docker registry is used to store images that are not meant to be shared publicly. You can set up your own private registry to maintain control over image distribution, security, and access within your organization.

Example: Using Docker Hub

1. **Pulling an Image:** Let's say you want to use the official `nginx` web server image from Docker Hub.

```
docker pull nginx
```



2. **Pushing an Image (Private Registry):** If you have your own private registry, you can push an image to it. Here, we'll use a hypothetical private registry at `registry.example.com`.

```
docker tag nginx registry.example.com/my-nginx  
docker push registry.example.com/my-nginx
```



3. **Pulling an Image from Private Registry:** To pull an image from your private registry, you need to log in first.

```
docker login registry.example.com  
docker pull registry.example.com/my-nginx
```



13. What is an entry point?

In Docker, the `ENTRYPOINT` instruction is used in a Dockerfile to specify the command that will be executed when a container is started from an image. It essentially sets the default command that will run as the primary process inside the container. This command can be overridden when the container is run by providing additional arguments.

Here's a brief explanation of the `ENTRYPOINT` instruction:

Usage: The `ENTRYPOINT` instruction is used in a Dockerfile to define the primary command to run when the container starts.

Example: Consider a Dockerfile for a Python application that runs a script named `app.py`. You can use the `ENTRYPOINT` instruction to set the default command for the container:

```
FROM python:3.9

WORKDIR /app

COPY app.py .

ENTRYPOINT ["python", "app.py"]
```

In this example, the `ENTRYPOINT` is set to `["python", "app.py"]`. When you run a container from this image, the `app.py` script will be executed as the primary process.

14. How to implement CI/CD in Docker?

Refer Blog : [How to build a CI/CD pipeline with Docker](#)

15. Will data on the container be lost when the docker container exits?

Yes, by default, any data stored within a Docker container will be lost when the container exits. Docker containers are designed to be ephemeral, meaning they are created and destroyed as needed. When a container is stopped or removed, the data within it, such as files, databases, and changes made during its runtime, will not persist unless you take specific steps to ensure data persistence.

To retain data across container restarts or removals, you can use the following approaches:

1. **Docker Volumes:** Docker volumes are a way to store data outside the container's filesystem, ensuring that the data persists even if the container is removed. Volumes can be used to store application data, configuration files, and databases. You can attach volumes to containers using the `-v` or `--volume` flag when running the `docker run` command.

Example:

```
docker run -v /path/on/host:/path/in/container my-app
```

2. **Docker Bind Mounts:** Similar to volumes, bind mounts allow you to map a host system directory into a container. Changes made to files in the bind mount on the host are reflected immediately in the container, and vice versa. This provides more flexibility compared to volumes.

Example:

```
docker run -v /host/directory:/container/directory my-app
```



3. **Named Volumes:** Docker also supports creating named volumes, which are managed by Docker and can be shared between multiple containers. Named volumes make it easier to manage and reuse data across containers.

Example:

```
docker run -v my-volume:/data my-app
```



4. **Docker Compose:** When using Docker Compose, you can define volumes in your `docker-compose.yml` file to ensure data persistence across services defined in the compose file.

Example:

```
version: "3"
services:
  web:
    image: nginx
    volumes:
      - my-volume:/usr/share/nginx/html
volumes:
  my-volume:
```



By using these methods, you can ensure that important data and changes made within a Docker container are preserved even after the container exits, restarts, or is removed.

16. What is a Docker swarm?

Docker Swarm, also known as Docker Swarm Mode, is a native clustering and orchestration solution for Docker. It allows you to create and manage a cluster of Docker nodes (machines) that work together as a single virtual Docker host. Docker Swarm provides a simple and integrated way to deploy and manage containerized applications at scale.

Here's a brief overview of Docker Swarm:

Key Concepts:

1. **Manager Nodes:** These are the control plane nodes that manage the cluster and orchestrate the deployment of services and tasks.
2. **Worker Nodes:** These are the worker machines that run containers based on the instructions provided by the manager nodes.
3. **Services:** Services are the declarative definitions of tasks to be run on the swarm. They define how many replicas of a container should run, which image to use, ports to expose, and more.

4. **Tasks:** Tasks are individual units of work that run containers as part of a service. Each replica of a service is a task.
5. **Load Balancing:** Docker Swarm provides built-in load balancing for services across the worker nodes.
6. **Overlay Networking:** Swarm supports overlay networking, allowing containers to communicate with each other across nodes as if they were on the same network.
7. **Rolling Updates and Rollbacks:** Swarm enables rolling updates and rollbacks for services, ensuring minimal downtime during updates.
8. **High Availability:** Swarm supports high availability by allowing you to create manager node replicas, preventing a single point of failure.

17. What are the docker commands for the following:

- **View Running Containers:**

```
docker ps
```



- **Run Container with a Specific Name:**

```
docker run --name my-container my-image
```



- **Export a Docker Container:**

```
docker export my-container > container-export.tar
```



- **Import an Existing Docker Image:**

```
docker import container-export.tar my-imported-image
```



- **Delete a Container:**

```
docker rm my-container
```



- **Remove Stopped Containers, Unused Networks, Build Caches, and Dangling Images:**

```
docker system prune
```



You can also use the `-a` flag to remove all unused images, not just dangling ones.

```
docker system prune -a
```



Parag Pallav Singh 

